

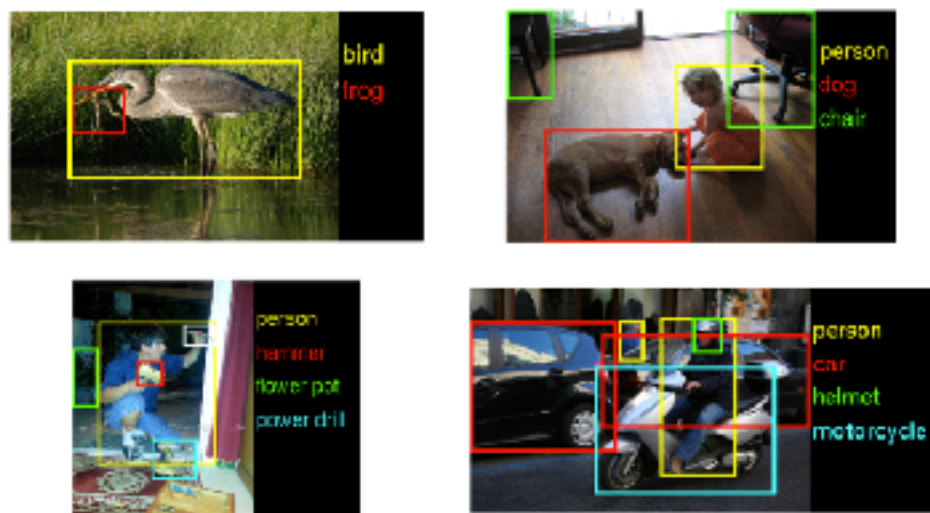
# Introduction to Deep Learning and Differentiable Physics

François Rousseau

# Why Deep Learning?

- Why is deep learning so popular?
  - Very good performance
  - End-to-end approach
- Why now?
  - Hardware: GPU
  - Datasets: ImageNet, Kaggle, etc.
  - User friendly libraries: PyTorch or Keras (and previously Caffe & Theano)
  - Algorithmics advances

# Large Scale Visual Recognition Challenge



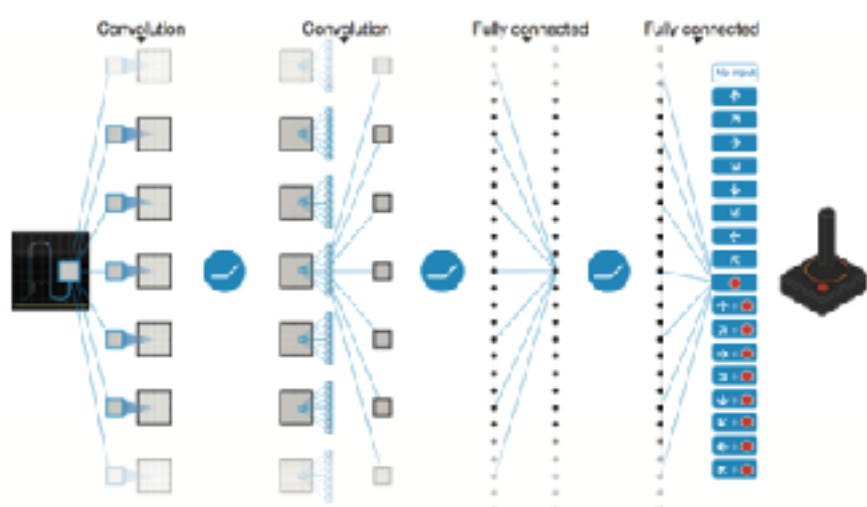
# Real Time Recognition




# Image captioning



# Algorithm playing ATARI



# Real Time Translation



# Text generator




# Overview

Machine  
Learning

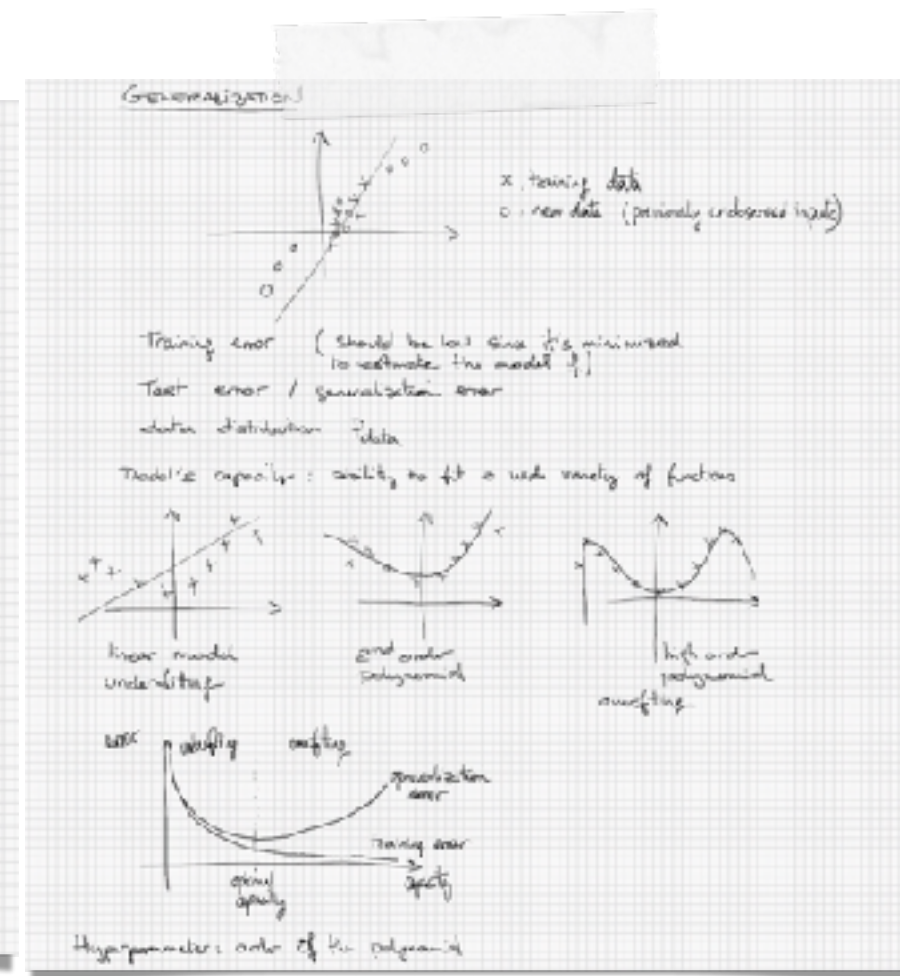
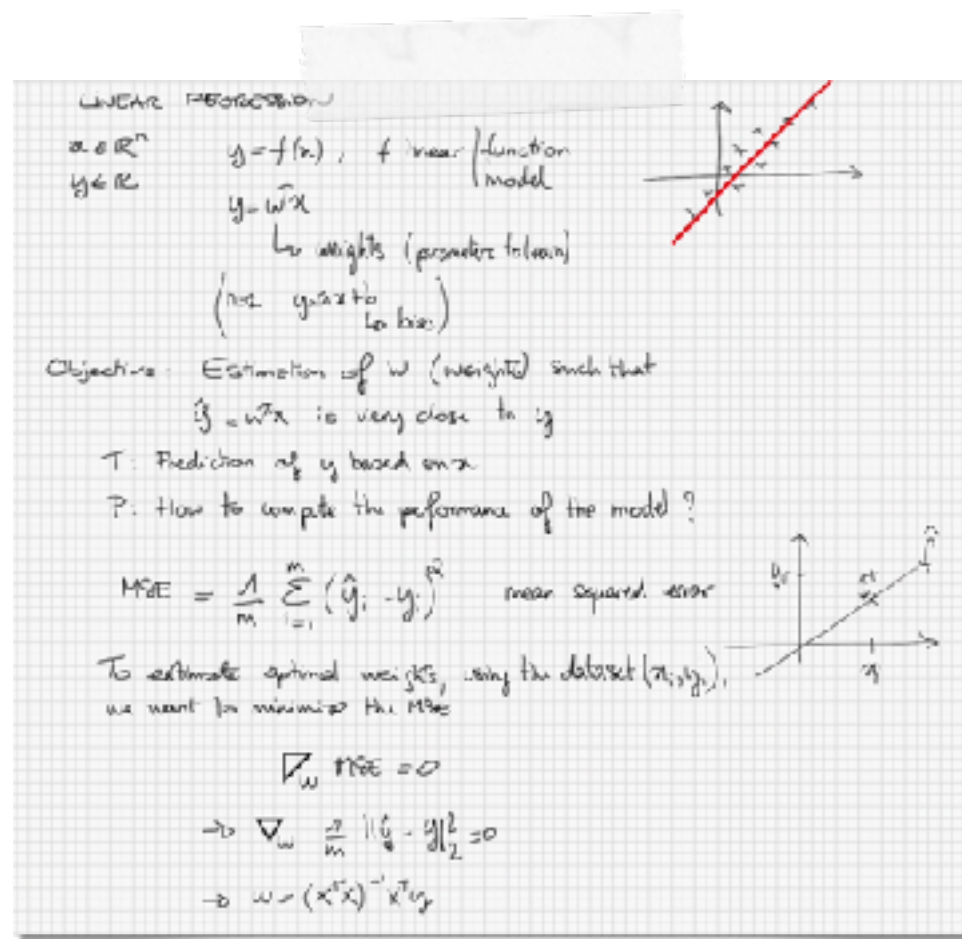
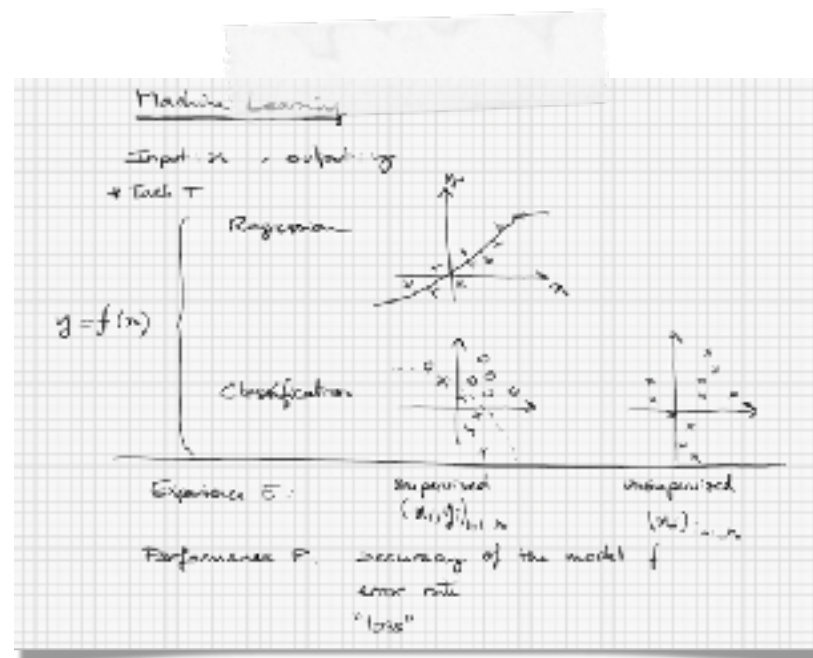
# Machine learning vocabulary

- **Data:** input  $x$  and output  $y$
- **Task:** regression, classification
- **Experience:** supervised  $(x_i, y_i)$  or unsupervised  $(x_i)$
- **Performance measure:** accuracy of the model or error rate





# Basics of Machine Learning




# Overview

Machine  
Learning

Neural Networks

# Neural networks


$$y = f(x; \theta)$$


Parameters to learn

- **Objective:** approximate a function  $f^*$
- **NN:** composition of functions  $f = f_n \circ \dots \circ f_2 \circ f_1$
- **How:** estimate parameters  $\theta$  of the neural network via a minimization problem:  $\arg \min_{\theta} \sum_i (f(x_i; \theta) - y_i)^2$

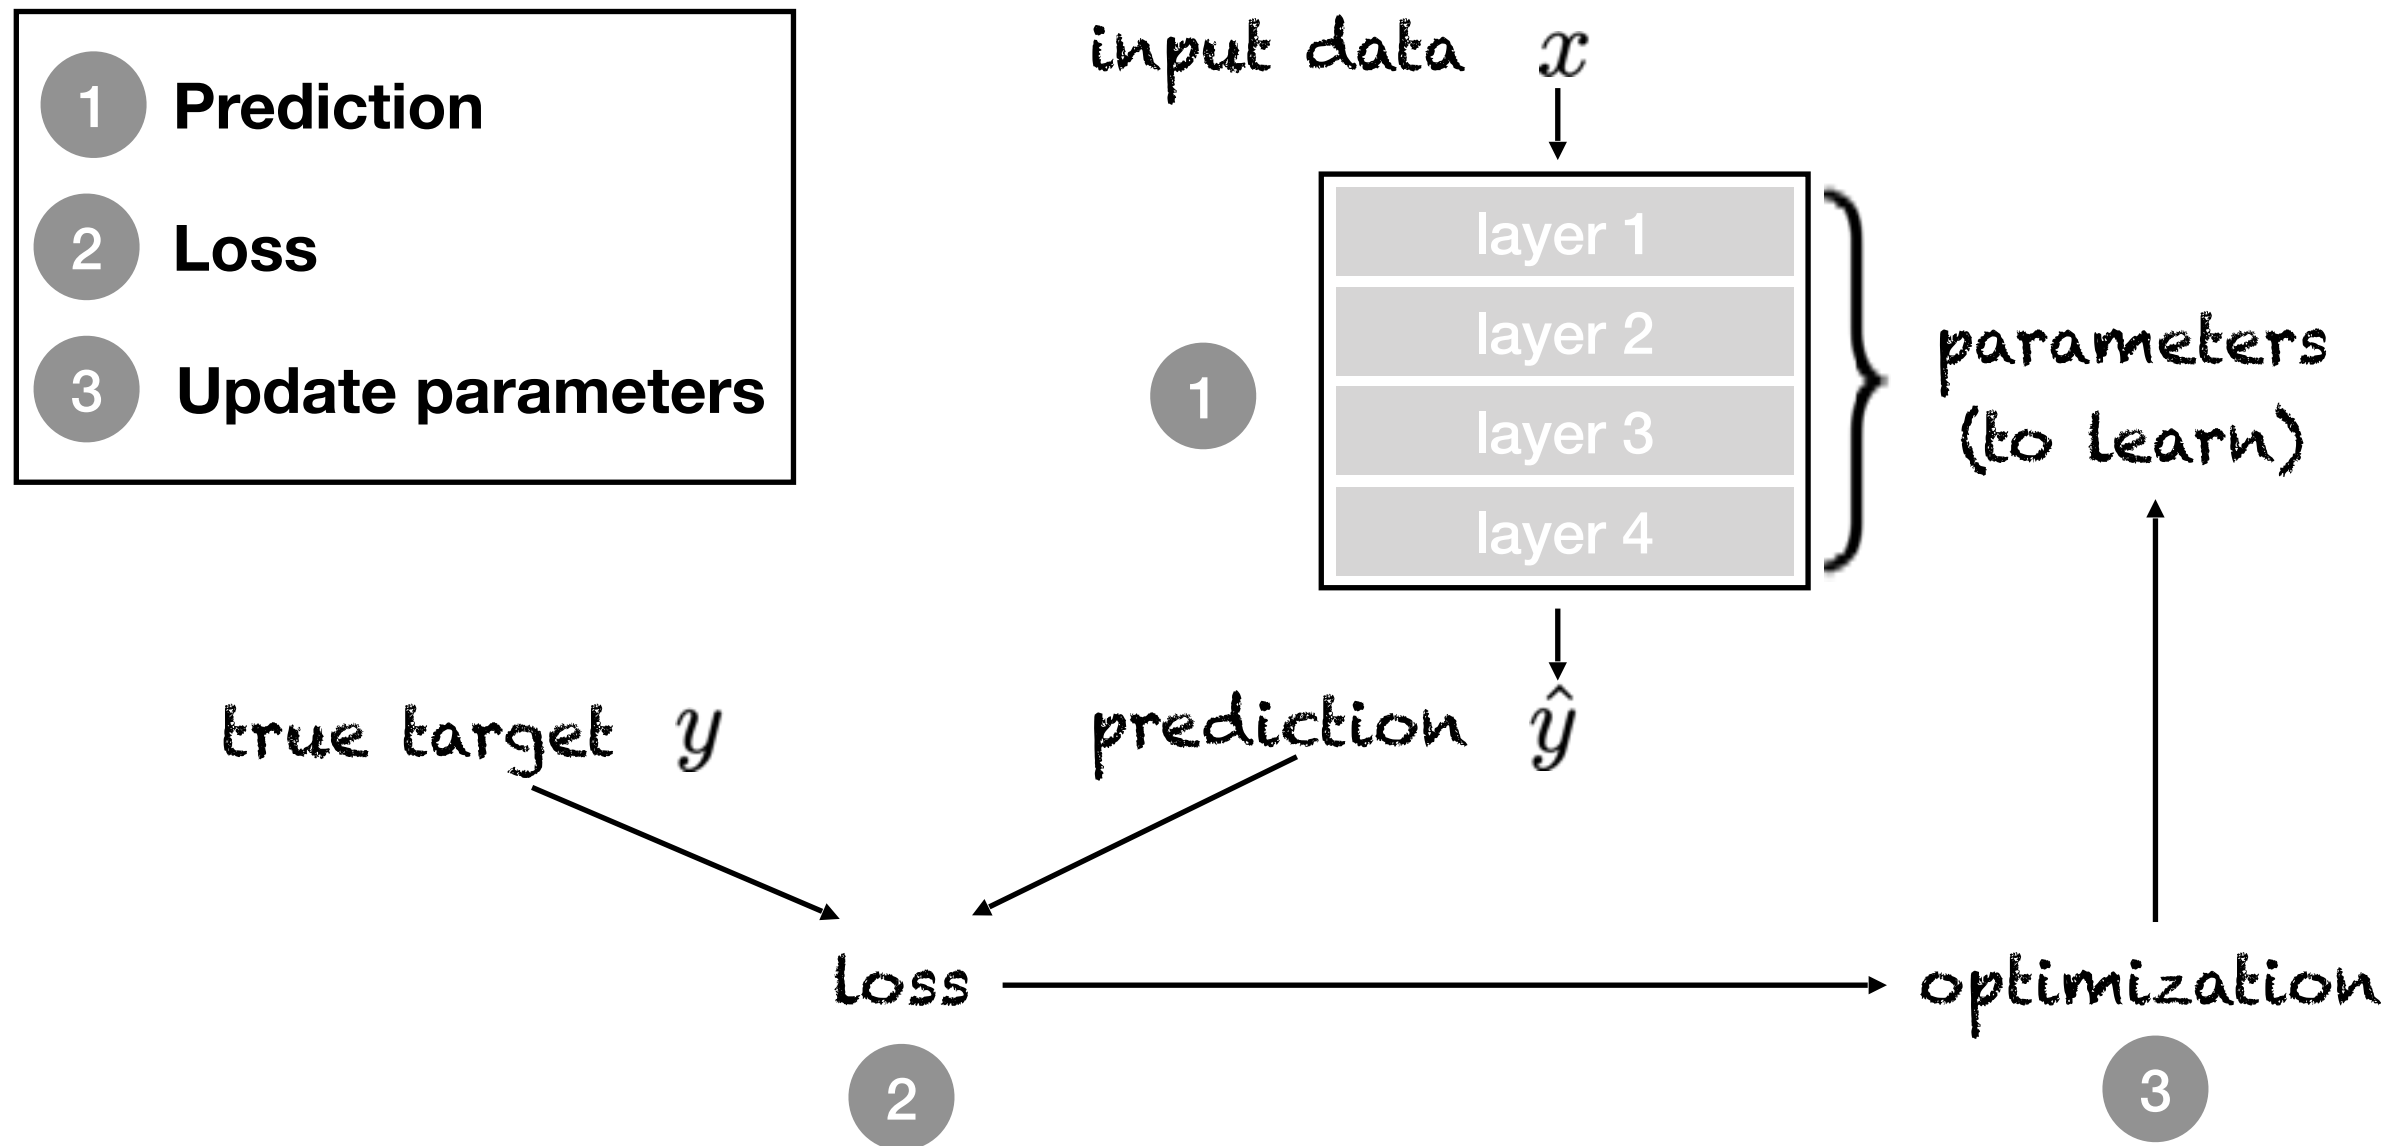


# Neural networks




$$f(x) = \sigma(\sum_j w_j \cdot h_j) = \sigma(\sum_j w_j \cdot \sigma(\sum_i w_{ij} x_i))$$

# Overview of NN



# In practice



```
import torch
import torch.nn as nn
```

```
class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(4,3)
        self.fc2 = nn.Linear(3,1)

    def forward(self, x):
        x = self.fc1(x)
        x = F.relu(x)
        x = self.fc2(x)
        x = F.softmax(x)
        return x
```


```
net = Net()
```

```
criterion = nn.MSELoss()
```

```
optimizer = torch.optim.SGD(net.parameters(), lr=0.01)
```

```
for i in range(epochs):          # training loop
    y_pred = net(x)              # compute the prediction using current parameters
    loss = criterion(y_pred,y)   # compute loss
    optimizer.zero_grad()        # initialize the gradient to zero
    loss.backward()              # back propagation
    optimizer.step()             # update the parameters values using the gradient
```

# Universal Approximation Theorem





1 hidden layer perceptron:

$$y \approx f(x, a, b, c) \stackrel{\text{def.}}{=} \sum_{i=1}^p c_i \varphi(a_i x + b_i)$$

Training:

$$\min_{a, b, c} E(a, b, c) \stackrel{\text{def.}}{=} \frac{1}{2n} \sum_{k=1}^n |y_k - f(x_k, a, b, c)|^2$$



[From Gabriel Peyré]

# Gradient-based approach for parameter estimation

- Gradient:

$$\frac{\partial J(\theta)}{\partial \theta} = \lim_{\delta \theta \rightarrow 0} \frac{J(\theta + \delta \theta) - J(\theta)}{\delta \theta}$$

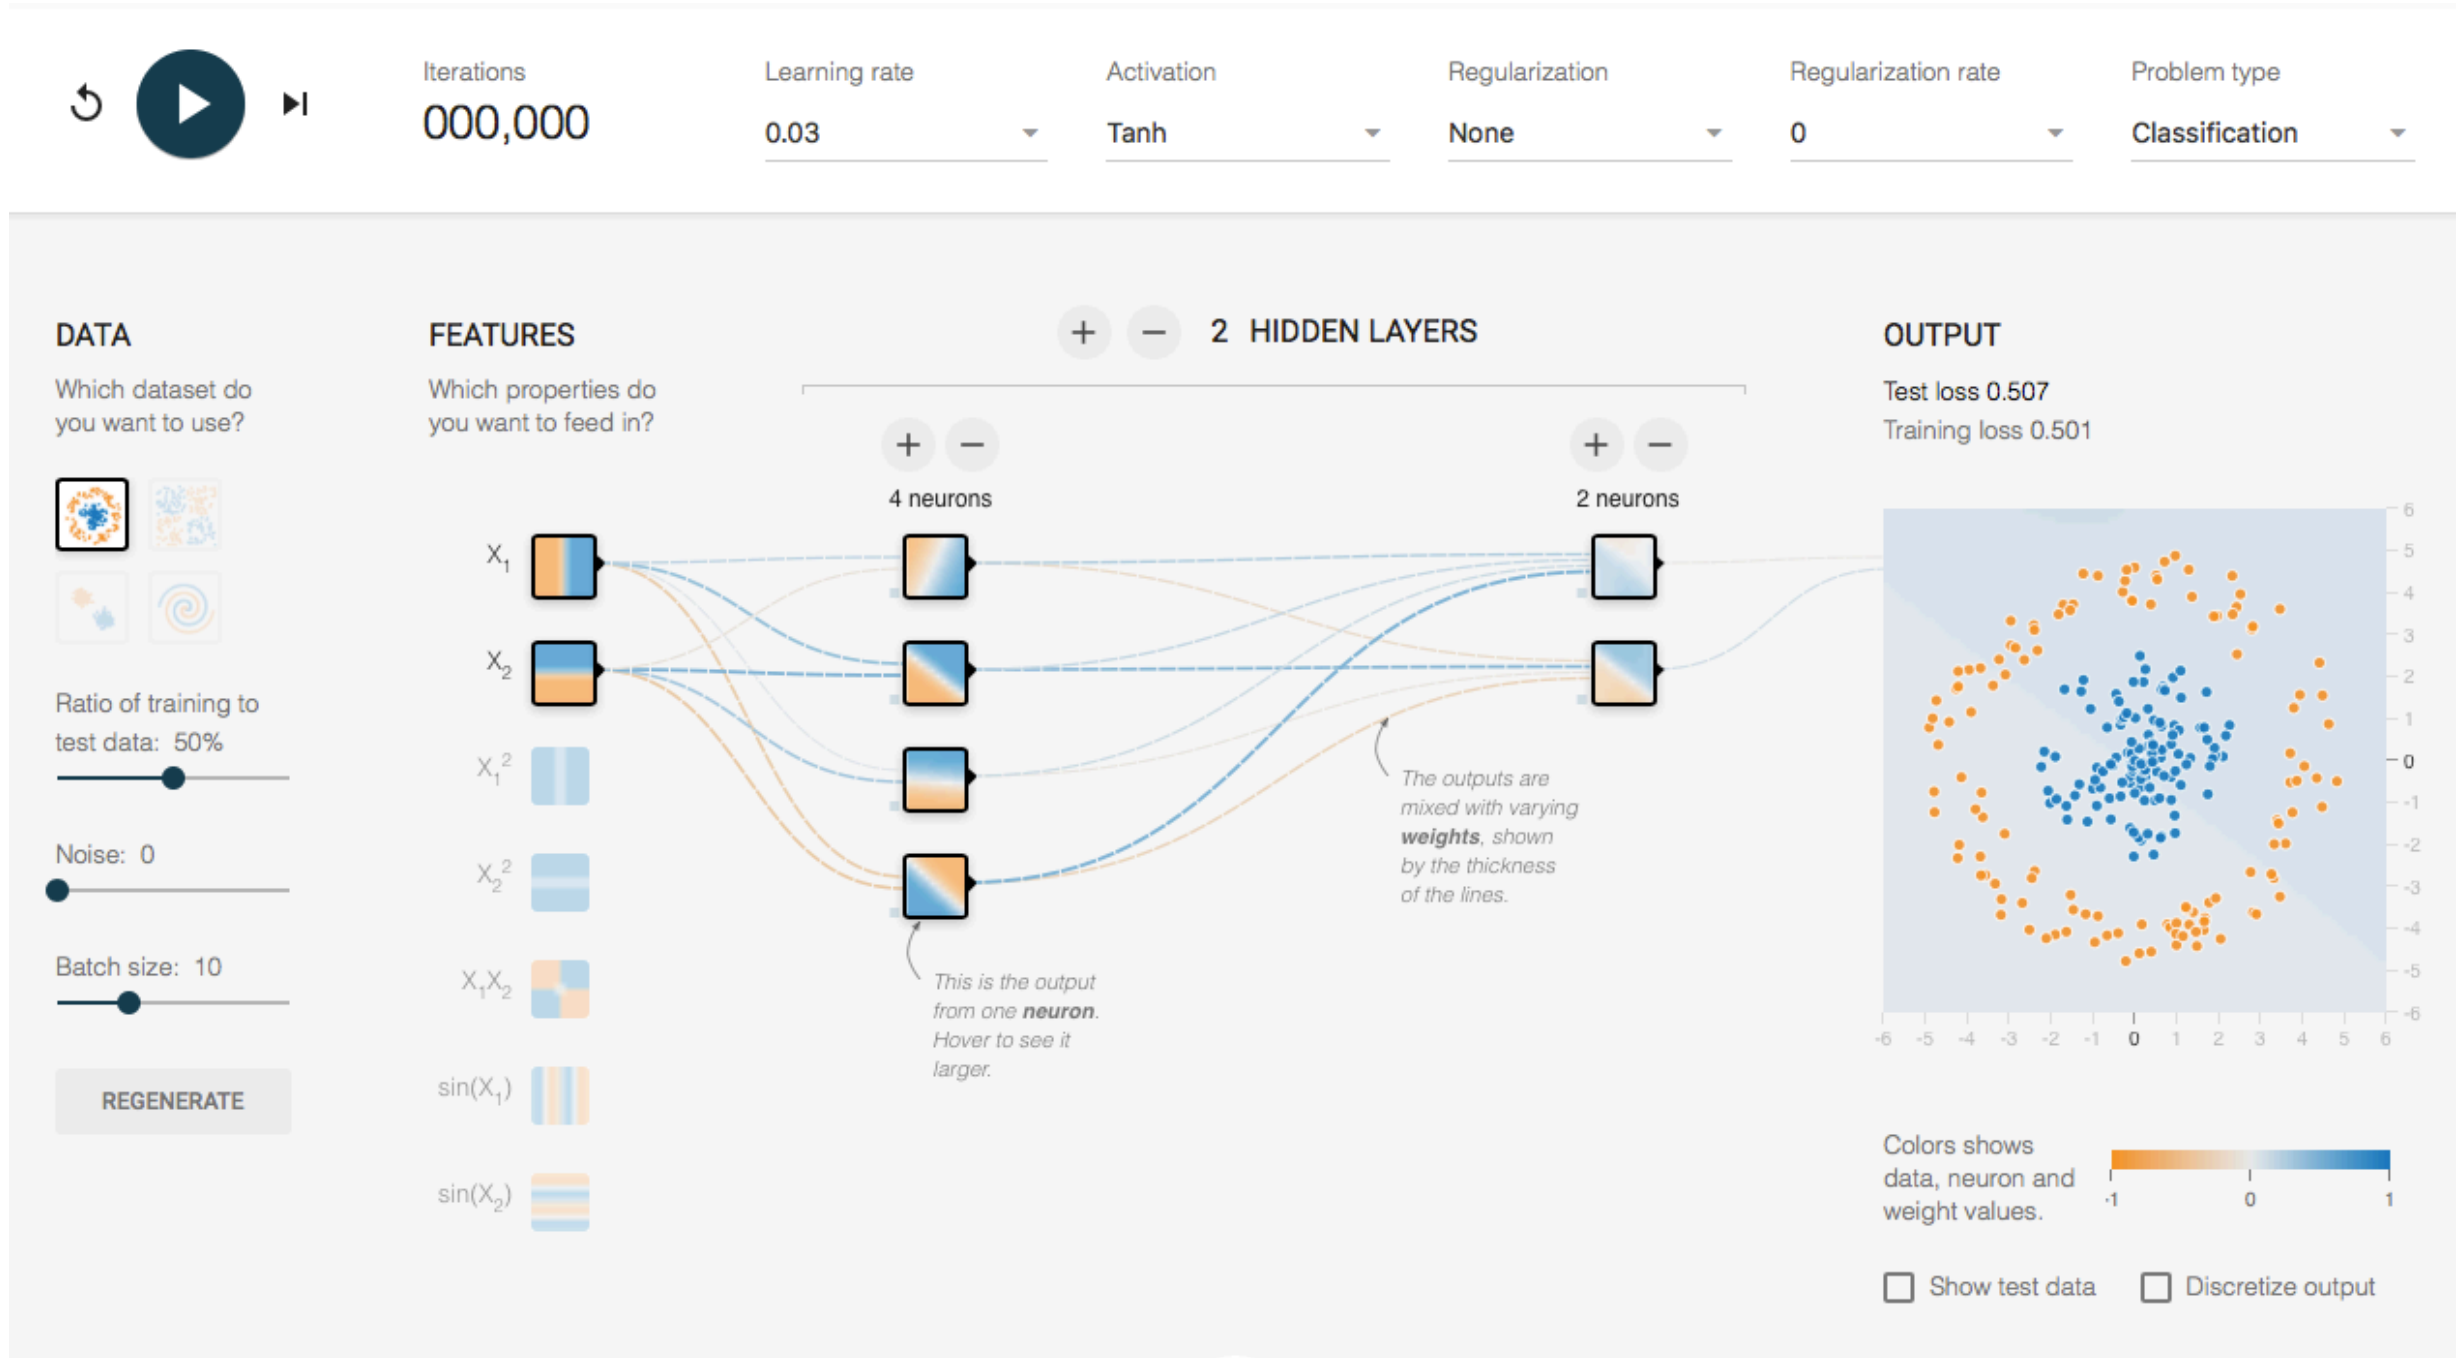
- Gradient descent:

$$\frac{\partial J(\theta)}{\partial \theta} = 0$$

$$\theta^{k+1} = \theta^k - \epsilon_k \frac{\partial J(\theta^k)}{\partial \theta^k}$$

Learning rate

# Practicing Deep Learning





Introduction\_regression.ipynb



# Overview

**Machine  
Learning**

**Neural Networks**

**Automatic  
Differentiation**

# Automatic differentiation

- How to compute the derivative of a function  $f$ ?

$$f'(a) = \frac{df}{dx}(a) = \lim_{h \rightarrow 0} \frac{f(a+h) - f(a)}{h}$$

- Unlike finite difference method, AD does not incur truncation errors.

# AD example

Given this:  $y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$

How to compute the value of  $y$  for  $x_1 = 1.5$  and  $x_2 = 0.5$ ?

$$v_{-1} = x_1 = 1.5000$$

$$v_0 = x_2 = 0.5000$$

# AD example

Given this:  $y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$

How to compute the value of  $y$  for  $x_1 = 1.5$  and  $x_2 = 0.5$ ?

$$\begin{array}{rclcl} v_{-1} & = & x_1 & = & 1.5000 \\ v_0 & = & x_2 & = & 0.5000 \\ \hline v_1 & = & v_{-1}/v_0 & = & 1.5000/0.5000 = 3.0000 \end{array}$$

# AD example

Given this:  $y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$


How to compute the value of  $y$  for  $x_1 = 1.5$  and  $x_2 = 0.5$ ?

$v_{-1}$	$=$	$x_1$	$=$	1.5000	
$v_0$	$=$	$x_2$	$=$	0.5000	
<hr/>					
$v_1$	$=$	$v_{-1}/v_0$	$=$	$1.5000/0.5000$	$=$ 3.0000
$v_2$	$=$	$\sin(v_1)$	$=$	$\sin(3.0000)$	$=$ 0.1411
$v_3$	$=$	$\exp(v_0)$	$=$	$\exp(0.5000)$	$=$ 1.6487
$v_4$	$=$	$v_1 - v_3$	$=$	$3.0000 - 1.6487$	$=$ 1.3513
$v_5$	$=$	$v_2 + v_4$	$=$	$0.1411 + 1.3513$	$=$ 1.4924
$v_6$	$=$	$v_5 * v_4$	$=$	$1.4924 * 1.3513$	$=$ 2.0167
<hr/>					
$y$	$=$	$v_6$	$=$	2.0167	

Evaluation trace

# AD example

Given this:  $y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$




Computational graph

# Forward mode

We want to differentiate the output variable  $y$  with respect to  $x_1$ .

$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$$



$$\dot{v}_i = \frac{\partial v_i}{\partial x_1}$$

$$\dot{v}_1 = \frac{\partial v_1}{\partial x_1} = \frac{\partial v_1}{\partial v_{-1}} \frac{\partial v_{-1}}{\partial x_1} + \frac{\partial v_1}{\partial v_0} \frac{\partial v_0}{\partial x_1}$$

$$\dot{v}_1 = \frac{1}{v_0} \dot{v}_{-1} + v_{-1} \left( -\frac{\dot{v}_0}{v_0^2} \right)$$

$$v_{-1} = x_1$$

$$v_0 = x_2$$

---



$$v_1 = v_{-1}/v_0$$



# Forward mode

$v_{-1} = x_1$	$= 1.5000$	
$\dot{v}_{-1} = \dot{x}_1$	$= 1.0000$	
$v_0 = x_2$	$= 0.5000$	
$\dot{v}_0 = \dot{x}_2$	$= 0.0000$	
<hr/>		
$v_1 = v_{-1}/v_0$	$= 1.5000/0.5000$	$= 3.0000$
$\dot{v}_1 = (\dot{v}_{-1} - v_1 * \dot{v}_0)/v_0$	$= 1.0000/0.5000$	$= 2.0000$


$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$$



# Forward mode

$v_{-1} = x_1$	$= 1.5000$	
$\dot{v}_{-1} = \dot{x}_1$	$= 1.0000$	
$v_0 = x_2$	$= 0.5000$	
$\dot{v}_0 = \dot{x}_2$	$= 0.0000$	
<hr/>		
$v_1 = v_{-1}/v_0$	$= 1.5000/0.5000$	$= 3.0000$
$\dot{v}_1 = (\dot{v}_{-1} - v_1 * \dot{v}_0)/v_0$	$= 1.0000/0.5000$	$= 2.0000$
$v_2 = \sin(v_1)$	$= \sin(3.0000)$	$= 0.1411$
$\dot{v}_2 = \cos(v_1) * \dot{v}_1$	$= -0.9900 * 2.0000$	$= -1.9800$


$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$$



# Forward mode

$v_{-1} = x_1$	$= 1.5000$	
$\dot{v}_{-1} = \dot{x}_1$	$= 1.0000$	
$v_0 = x_2$	$= 0.5000$	
$\dot{v}_0 = \dot{x}_2$	$= 0.0000$	
<hr/>		
$v_1 = v_{-1}/v_0$	$= 1.5000/0.5000$	$= 3.0000$
$\dot{v}_1 = (\dot{v}_{-1} - v_1 * \dot{v}_0)/v_0$	$= 1.0000/0.5000$	$= 2.0000$
$v_2 = \sin(v_1)$	$= \sin(3.0000)$	$= 0.1411$
$\dot{v}_2 = \cos(v_1) * \dot{v}_1$	$= -0.9900 * 2.0000$	$= -1.9800$
$v_3 = \exp(v_0)$	$= \exp(0.5000)$	$= 1.6487$
$\dot{v}_3 = v_3 * \dot{v}_0$	$= 1.6487 * 0.0000$	$= 0.0000$
$v_4 = v_1 - v_3$	$= 3.0000 - 1.6487$	$= 1.3513$
$\dot{v}_4 = \dot{v}_1 - \dot{v}_3$	$= 2.0000 - 0.0000$	$= 2.0000$
$v_5 = v_2 + v_4$	$= 0.1411 + 1.3513$	$= 1.4924$
$\dot{v}_5 = \dot{v}_2 + \dot{v}_4$	$= -1.9800 + 2.0000$	$= 0.0200$
$v_6 = v_5 * v_4$	$= 1.4924 * 1.3513$	$= 2.0167$
$\dot{v}_6 = \dot{v}_5 * v_4 + v_5 * \dot{v}_4$	$= 0.0200 * 1.3513 + 1.4924 * 2.0000$	$= 3.0118$
<hr/>		
$y = v_6$	$= 2.0100$	
$\dot{y} = \dot{v}_6$	$= 3.0110$	

$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$$



Note: need to do the same to compute  $\frac{\partial y}{\partial x_2}$

# Reverse mode

We want to calculate the sensitivity of an output variable with respect to each of the intermediate variables.


$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$$

$$\bar{v}_i = \frac{\partial y}{\partial v_i}$$

$$\bar{v}_6 = 1$$

$$\bar{v}_5 = \frac{\partial y}{\partial v_5} = \frac{\partial y}{\partial v_6} \frac{\partial v_6}{\partial v_5} = \bar{v}_6 v_4$$

$$\bar{v}_4 = \frac{\partial y}{\partial v_4} = \frac{\partial y}{\partial v_6} \frac{\partial v_6}{\partial v_4} + \frac{\partial y}{\partial v_5} \frac{\partial v_5}{\partial v_4} = \bar{v}_6 v_5 + \bar{v}_5 \times 1$$



$$v_4 = v_1 - v_3$$

$$v_5 = v_2 + v_4$$

$$v_6 = v_5 * v_4$$

---


$$y = v_6$$

# Reverse mode

$$v_{-1} = x_1 = 1.5000$$

$$v_0 = x_2 = 0.5000$$

$$v_1 = v_{-1}/v_0 = 1.5000/0.5000 = 3.0000$$

$$v_2 = \sin(v_1) = \sin(3.0000) = 0.1411$$

$$v_3 = \exp(v_0) = \exp(0.5000) = 1.6487$$

$$v_4 = v_1 - v_3 = 3.0000 - 1.6487 = 1.3513$$


$$v_5 = v_2 + v_4 = 0.1411 + 1.3513 = 1.4924$$

$$v_6 = v_5 * v_4 = 1.4924 * 1.3513 = 2.0167$$

$$y = v_6 = 2.0167$$

$$\bar{v}_6 = \bar{y} = 1.0000$$

$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$$



# Reverse mode

$$v_{-1} = x_1 = 1.5000$$

$$v_0 = x_2 = 0.5000$$

$$v_1 = v_{-1}/v_0 = 1.5000/0.5000 = 3.0000$$

$$v_2 = \sin(v_1) = \sin(3.0000) = 0.1411$$

$$v_3 = \exp(v_0) = \exp(0.5000) = 1.6487$$

$$v_4 = v_1 - v_3 = 3.0000 - 1.6487 = 1.3513$$

$$v_5 = v_2 + v_4 = 0.1411 + 1.3513 = 1.4924$$


$$v_6 = v_5 * v_4 = 1.4924 * 1.3513 = 2.0167$$

$$y = v_6 = 2.0167$$

$$\bar{v}_6 = \bar{y} = 1.0000$$

$$\bar{v}_5 = \bar{v}_6 * v_4 = 1.0000 * 1.3513 = 1.3513$$

$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$$



# Reverse mode

$$v_{-1} = x_1 = 1.5000$$

$$v_0 = x_2 = 0.5000$$

$$v_1 = v_{-1}/v_0 = 1.5000/0.5000 = 3.0000$$

$$v_2 = \sin(v_1) = \sin(3.0000) = 0.1411$$

$$v_3 = \exp(v_0) = \exp(0.5000) = 1.6487$$

$$v_4 = v_1 - v_3 = 3.0000 - 1.6487 = 1.3513$$

$$v_5 = v_2 + v_4 = 0.1411 + 1.3513 = 1.4924$$

$$v_6 = v_5 * v_4 = 1.4924 * 1.3513 = 2.0167$$

$$y = v_6 = 2.0167$$

$$\bar{v}_6 = \bar{y} = 1.0000$$

$$\bar{v}_5 = \bar{v}_6 * v_4 = 1.0000 * 1.3513 = 1.3513$$

$$\bar{v}_4 = \bar{v}_6 * v_5 = 1.0000 * 1.4924 = 1.4924$$

$$\bar{v}_4 = \bar{v}_4 + \bar{v}_5 = 1.4924 + 1.3513 = 2.8437$$

$$\bar{v}_2 = \bar{v}_5 = 1.3513$$

$$\bar{v}_3 = -\bar{v}_4 = -2.8437$$

$$\bar{v}_1 = \bar{v}_4 = 2.8437$$

$$\bar{v}_0 = \bar{v}_3 * v_3 = -2.8437 * 1.6487 = -4.6884$$

$$\bar{v}_1 = \bar{v}_1 + \bar{v}_2 * \cos(v_1) = 2.8437 + 1.3513 * (-0.9900) = 1.5059$$


$$\bar{v}_0 = \bar{v}_0 - \bar{v}_1 * v_1/v_0 = -4.6884 - 1.5059 * 3.000/0.5000 = -13.7239$$

$$\bar{v}_{-1} = \bar{v}_1/v_0 = 1.5059/0.5000 = 3.0118$$

$$\bar{x}_2 = \bar{v}_0 = -13.7239$$

$$\bar{x}_1 = \bar{v}_{-1} = 3.0118$$

$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)] * [x_1/x_2 - \exp(x_2)]$$



Note: we have computed  $\frac{\partial y}{\partial x_1}$  and  $\frac{\partial y}{\partial x_2}$  but the computational graph is required.



# AD and jacobian matrix

$$F : \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix} \mapsto \begin{pmatrix} f_1(x_1, \dots, x_n) \\ \vdots \\ f_m(x_1, \dots, x_n) \end{pmatrix}$$

$$J_F(M) = \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \dots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \dots & \frac{\partial f_m}{\partial x_n} \end{pmatrix}$$

Forward pass: computation of a column

Reverse pass: computation of a row

In machine learning,  $m = 1$  (scalar value of the loss function), so we use the reverse mode.

# Back-propagation

- Back-propagation is a special case of reverse mode automatic differentiation.
- Back-propagation is an efficient way to recursively compute the gradient:

$$\frac{\partial L}{\partial u} = \sum_i \frac{\partial L}{\partial v_i} \frac{\partial v_i}{\partial u}$$

parent of node u


# Back-propagation

$$Loss = \frac{1}{2} \sum_k (f(x_k) - y_k)^2$$

$$f(x_k) = \sigma(in_k)$$

$$\frac{\partial Loss}{\partial w_{jk}} = \frac{\partial Loss}{\partial in_k} \frac{\partial in_k}{\partial w_{jk}} = \delta_k \frac{\partial (\sum_j w_{jk} h_j)}{\partial w_{jk}} = \delta_k h_j$$

$$\frac{\partial Loss}{\partial w_{ij}} = \frac{\partial Loss}{\partial in_j} \frac{\partial in_j}{\partial w_{ij}} = \delta_j \frac{\partial (\sum_i w_{ij} x_i)}{\partial w_{ij}} = \delta_j x_i$$



$$\delta_k = \frac{\partial}{\partial in_k} \left( \sum_k \frac{1}{2} [\sigma(in_k) - y_k]^2 \right) = [\sigma(in_k) - y_k] \sigma'(in_k)$$

$$\delta_j = \sum_k \frac{\partial Loss}{\partial in_k} \frac{\partial in_k}{\partial in_j} = \sum_k \delta_k \cdot \frac{\partial}{\partial in_j} \left( \sum_j w_{jk} \sigma(in_j) \right) = [\sum_k \delta_k w_{jk}] \sigma'(in_j)$$

# Back-propagation

$$Loss = \frac{1}{2} \sum_k (f(x_k) - y_k)^2$$
$$f(x_k) = \sigma(in_k)$$

Stochastic Gradient Descent:


$$\theta^{k+1} = \theta^k - \epsilon_k \frac{\partial J(\theta^k)}{\partial \theta^k}$$

$$\frac{\partial Loss}{\partial w_{jk}} = \delta_k h_j = [\sigma(in_k) - y_k] \sigma'(in_k) h_j$$

$$\frac{\partial Loss}{\partial w_{ij}} = \delta_j x_i = [\sum_k \delta_k w_{jk}] \sigma'(in_j) x_i$$

Updates involve **scaled error from output** and **input feature**

After forward pass, compute  $\delta_k$  from final layer and then  $\delta_j$  for previous layer.



# Overview

Machine  
Learning

Neural Networks


Automatic  
Differentiation

Supervised  
Learning



Automatic\_differentiation.ipynb

# About supervised learning

$$y = f(x; \theta)$$


Parameters to learn

- Supervised learning is a natural starting point for any DL project
- Supervised training is stable (and provide an upper bound)
- Garbage-in-garbage-out: know your data!
- Supervised DL is interpolation (generalization issue)
- Fully connected layers (MLPs) vs CNN



# Overview

Machine  
Learning

Neural Networks


Automatic  
Differentiation

Supervised  
Learning

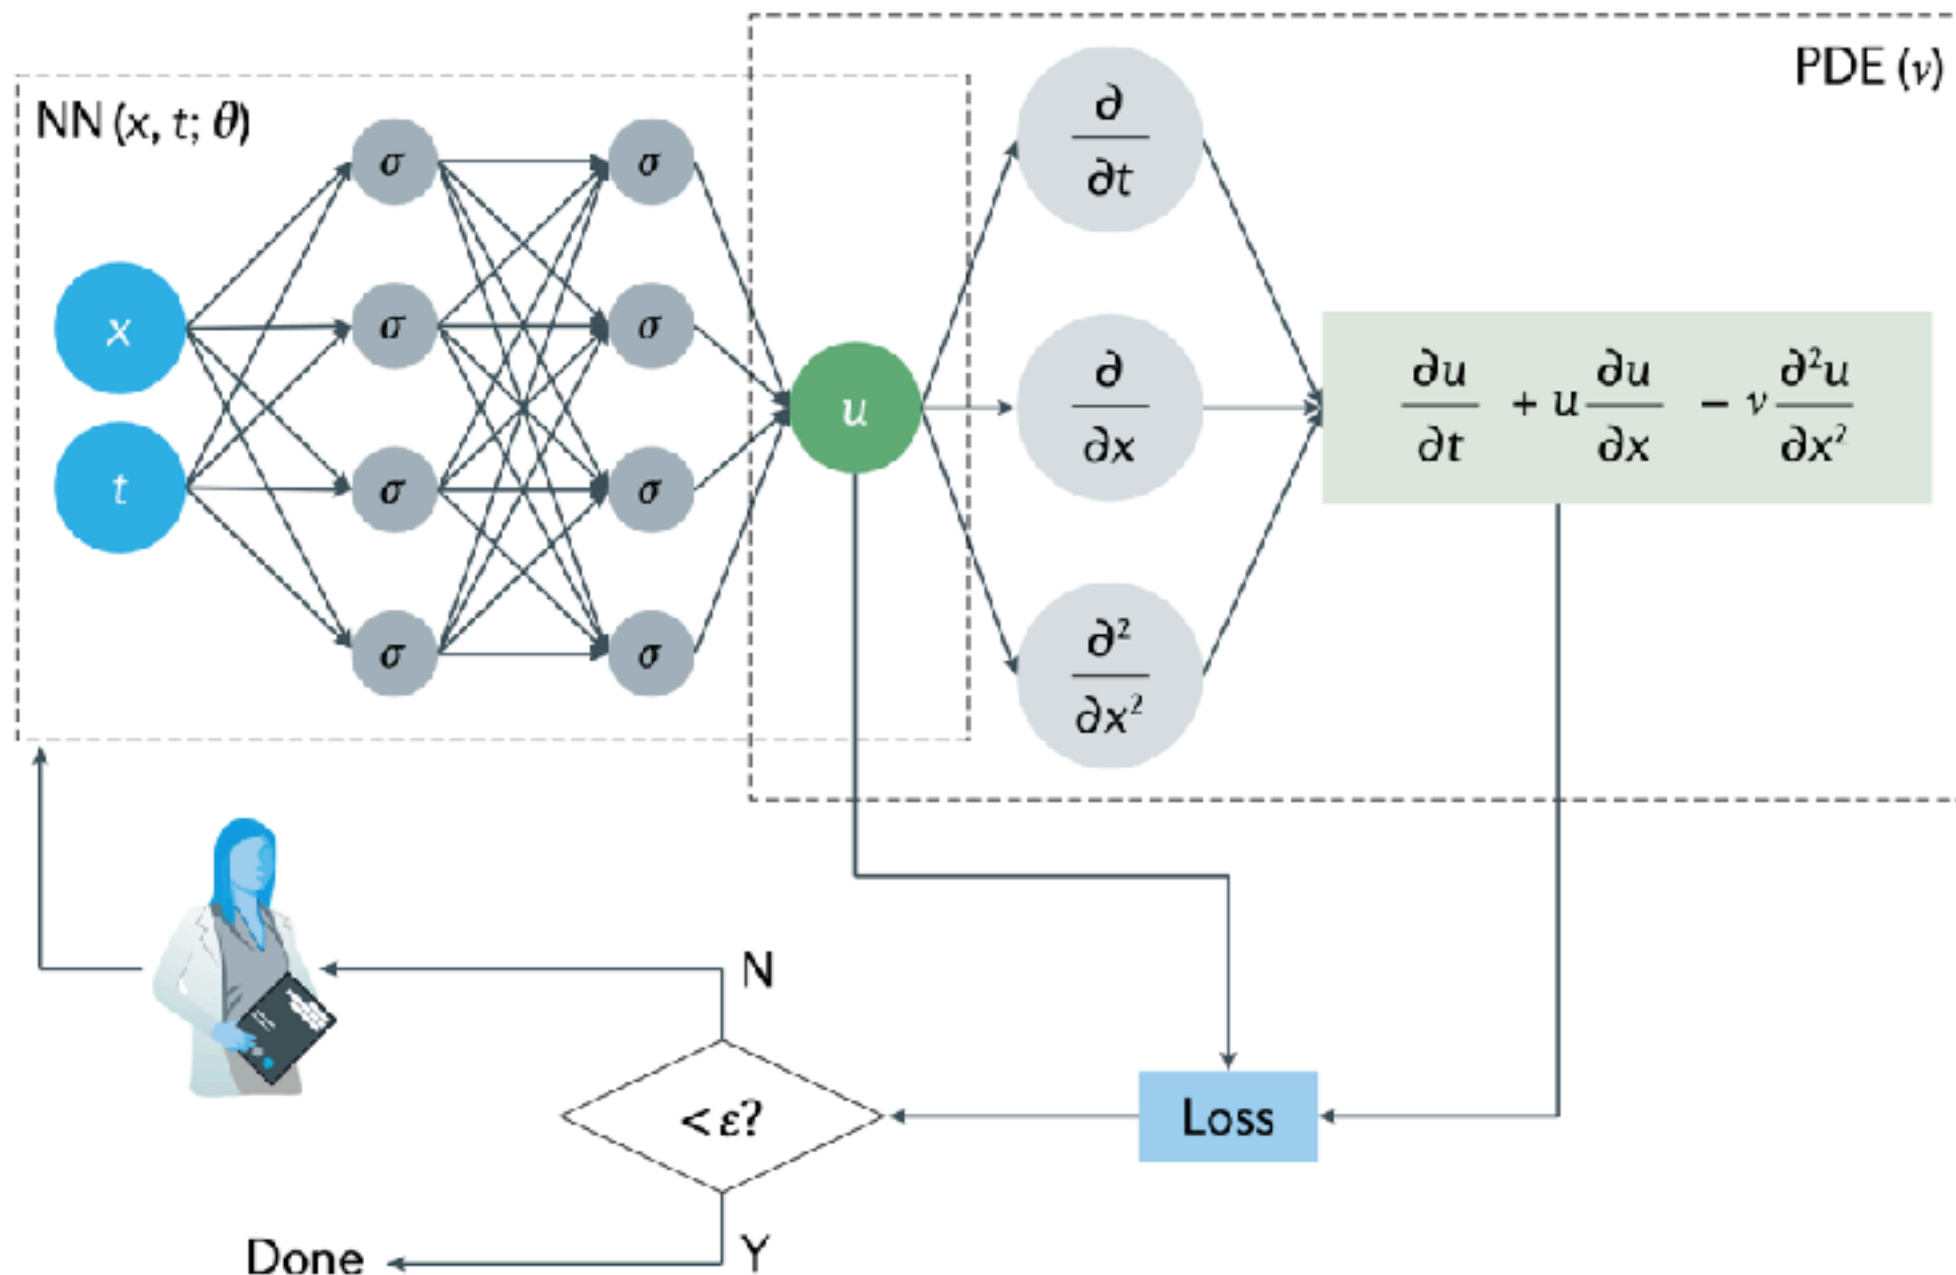
Physics-  
informed  
Learning

# Burgers' equation

$$\frac{\partial u}{\partial t} + u \nabla u = \nu \nabla \cdot \nabla u$$



# Physics-informed NN



# Physics-informed NN

- Given a PDE for  $\mathbf{u}(\mathbf{x}, t)$ :  $\mathbf{u}_t = \mathcal{F}(\mathbf{u}_x, \mathbf{u}_{xx}, \dots)$
- To approximate  $\mathbf{u}$  with a neural network, we want to minimize a residual term:  $R = \mathbf{u}_t - \mathcal{F}(\mathbf{u}_x, \mathbf{u}_{xx}, \dots)$
- In a supervised setting, the training objective becomes:

$$\arg \min_{\theta} \sum_i \alpha_0 (f(x_i; \theta) - y_i)^2 + \alpha_1 R(f(x_i; \theta))$$



Burgers\_1D.ipynb

# About Physics-informed NN

- Physical equations are included in the form of soft constraints
- Automatic differentiation can be use to compute high-order derivatives in PDEs
- Each derivative computation can be memory expensive (and slow)
- PINN is a inverse problem with physical regularization

# Overview

Machine  
Learning

Neural Networks

Automatic  
Differentiation

Supervised  
Learning

Physics-  
informed  
Learning

Differentiable  
physics

# Differentiable physics

- Consider a continuous formulation  $\mathcal{P}(\mathbf{x}, \nu)$  of a physical quantity of interest  $\mathbf{u}(\mathbf{x}, t) : \mathbb{R}^d \times \mathbb{R}^+ \rightarrow \mathbb{R}^d$ .
- Assume that  $\mathcal{P}(\mathbf{x}, \nu)$  can be decomposed as a sequence of operations  $\mathcal{P}_1, \mathcal{P}_2, \dots, \mathcal{P}_m$  such that:

$$\mathbf{u}(t + \Delta t) = \mathcal{P}_1 \circ \mathcal{P}_2 \circ \dots \circ \mathcal{P}_m(\mathbf{u}(t), \nu)$$

- Objective: find  $\mathbf{u}$  that minimizes a loss given hard physical constraints (given by  $\mathcal{P}$ )



# Differentiable physics

- Physical model  $\mathcal{P}$ : 
$$\frac{\partial d}{\partial t} + \mathbf{u} \nabla d = 0$$
- Evolution equation:  $d(t + \Delta t) = \mathcal{P}(d(t), \mathbf{u}, t + \Delta t)$
- Objective: find a motion  $\mathbf{u}$  that transform a given initial state  $d^0$  to a target density  $d^{target}$  at final time  $t^e$
- Minimization problem:

$$\arg \min_{\mathbf{u}} \|\mathcal{P}(d^0, \mathbf{u}, t^e) - d^{target}\|^2$$

# Burgers' equation and Differentiable physics

## PhiFlow code

```
LR = 5.



grads=[]
for optim_step in range(5):
    velocities = [velocity]
    with math.record_gradients(velocity.values):
        for time_step in range(STEPS):
            v1 = diffuse.explicit(1.0*velocities[-1], NU, DT)
            v2 = advect.semi_lagrangian(v1, v1, DT)
            velocities.append(v2)

        loss = field.l2_loss(velocities[16] - SOLUTION_T16)*2./N # MSE
        print('Optimization step %d, loss: %f' % (optim_step,loss))

        grads.append( math.gradients(loss, velocity.values) )

velocity = velocity - LR * grads[-1]
```

# Burgers' equation and Differentiable physics




# About Differentiable physics

- Makes use of physical model and numerical methods for discretization
- Efficient evaluation of simulation and derivatives
- More complicated implementation (than PINN)


# Take home messages

- Machine learning: data driven approaches
- Neural networks: composition of functions  $f = f_n \circ \dots \circ f_2 \circ f_1$
- Automatic differentiation: powerful tool for gradient computation
- Supervised learning for regression: interpolation
- PINN : add physical loss term (regularization term)
- Differentiable physics: unfold the physical equation

# References



<https://www.deeplearningbook.org>



<https://d2l.ai>



<https://www.physicsbaseddeeplearning.org>